

A Network and Its Reflection: What True, False, and Noise Self-Model Broadcasts Do and Do Not Change in an Adaptive Multi-Agent Network

Ivan Bhargava and Harsh Modi

IIIT-Bangalore

`ivan.bhargava@iiitb.ac.in`, `harsh.modi@iiitb.ac.in`

WWW home page: <https://github.com/Ivan825/Project-ONE>

Abstract. Networks often act on measurements of their own state: descriptions can change what they describe. We build a controlled adaptive-network model in which node states and topology coevolve, and ask when feedback from global network observables causes genuine self-fulfilling effects. Across 3,127 deterministic runs with paired-seed comparisons, measurement without broadcast changes nothing. A seemingly strong “pull” toward false and replayed self-models disappears when untreated twin trajectories are scored against matched references: intrinsic dynamics mimic performativity. Broadcasts have real effects, reshaping cooperation and network structure, and whether a false self-model comes true, in the outcome sense, depends on how agents respond. Attention to the broadcast is heritable and evolves downward under consequential feedback — most steeply where that feedback carries high corrective drive — surviving three architectural controls and an ecological ensemble. These results distinguish apparent performativity from genuine feedback effects and show that reliance on a collective self-model can evolve.

Keywords: adaptive networks, network dynamics, agent-based modeling, reflexivity, performative prediction, evolution of trust

1 Introduction

Adaptive networks couple dynamics *on* the network — changing node states — with dynamics *of* the network — changing topology [1, 2]. A distinctive subclass has received far less controlled attention: networks that respond to *descriptions of themselves*. Financial markets move on published indices computed from their own trades [3]; universities restructure around rankings derived from their own behavior [4]; recommendation platforms amplify items *because* an internal measurement declared them trending. In each case a macroscopic measurement is compressed, published, and re-enters as an input to microscopic decisions — the loop Merton called the self-fulfilling prophecy [5], whose measure-distorting variant is Goodhart’s law [6].

Empirical study of this loop is confounded: one can rarely withhold the measurement, falsify it in controlled directions, or replay history under identical randomness. We therefore built a minimal simulated world where the loop




Fig. 1. The reflexive loop and its corruption points: local actions form the network, the observer compresses it into a self-model, and the channel returns it truthfully, distorted, replayed, as noise, or not at all. Propensity responses scale with $\gamma_i g$; hub-targeted pruning with γ_i alone.

can be opened, closed, and corrupted at will, every run is exactly reproducible from (configuration, seed), and agents are transparent trait vectors rather than learned policies, so every pathway from broadcast to response is inspectable.

Our contributions are: (1) an open, deterministic experimental platform coupling an evolving agent ecology, an adaptive interaction network, a global observer, and a controllable broadcast channel; (2) a six-condition causal design — verifying observer-side-effect invariance and isolating the effect of broadcasting, with matched-range noise and replayed-self-model controls — and a *passive-counterfactual test* for reflexivity metrics, showing that a natural “story pull” statistic can be fully explained by untreated twin trajectories; (3) evidence for what broadcasts *do* change: truthful feedback doubles costly cooperation, distorted broadcasts mobilize or pacify depending on the lie–response pairing (reversing under conformity), and false and noise broadcasts densify the network while also reducing hub concentration; and (4) receiver sensitivity as a continuous heritable trait whose *population mean changes systematically under feedback*, tracks the broadcast-implied corrective-drive intensity, and remains robust across three architectural controls and a randomized ecological ensemble — reliance on the observer is not updated by rule; it evolves. Figure 1 summarizes the design.

2 Related Work

Adaptive networks and global coupling. Coevolution of state and topology [1, 2] covers our setting; robustness to targeted node removal is classical [7], and we use hub removal as a standardized perturbation. A mean field fed back to units reshapes collective dynamics [8]; ours is multidimensional, potentially *false*, and its units are born, die, and evolve. **Reflexivity and performative prediction.** Self-fulfilling prophecy [5], reflexivity in markets [3], and the reactivity of rankings and metrics [4, 6] motivate our false-broadcast conditions; here they are dosable and ablatability. The formal baseline we build on is performative prediction — predictions that alter the distribution they forecast [19, 23], including networked settings [20], endogenous self-fulfilling signals [24], and — closest to us — populations deciding whether to trust predictions of collective outcomes [21]. That line updates trust from past accuracy and leaves causal identification without randomization open [22]; we add two elements: a passive-twin design exposing a concrete false-positive failure mode of directional-alignment metrics (Sect. 5.2), and a heritable receiver-sensitivity trait — inherited and

mutated with no rule referencing reliability — plus demographic turnover, absent from these models. **Signaling, trust, and agent-based modeling.** In sender–receiver games receivers can evolve to ignore unreliable signals [9]; here the signal is *self-referential* and attention to it is a continuous heritable trait in an energy economy, not a payoff matrix. Collective opinion and cooperation are ABM mainstays [16, 17]; we follow the generative tradition [10, 11]; the engine draws on Mesa [12] and NetworkX [13]. Multi-agent RL treats global state as an *architectural* aid [14]; here it is an *experimental variable* with controlled veracity. For language-model agent populations [15], this rule-based setting gives causally clean baselines for their natural shared-dashboard designs.

3 The Model

3.1 Agents and ecology

Each agent i carries a heritable trait vector $\theta_i = (\text{coop}, \text{expl}, \text{risk}, \text{soc}, \text{share}, \gamma_i) \in [0, 1]^6$, where γ_i is its *global sensitivity* — how strongly broadcasts modulate its behavior — and $\bar{\gamma}$ its population mean. Agents hold energy; each step they pay metabolism and choose one action — harvest, share (costly transfer to a neighbor), connect, prune, reproduce, or idle — with probabilities proportional to trait- and state-dependent propensities. Agents are the vertices of a time-varying undirected graph $G(t) = (V(t), E(t))$: *connect* and reproduction add interaction edges, *prune* and death remove them — topology coevolves with state, an adaptive/coevolutionary complex network [1] (Sect. 3.4). The resource pool regenerates logistically; reproduction requires an energy threshold and copies traits with Gaussian mutation ($\sigma = 0.05$, clipped to $[0, 1]$); energy exhaustion or age causes death; founding ages are staggered to avoid cohort artifacts.

3.2 Observer and broadcast conditions

Every $\Delta t = 10$ steps the observer computes the macrostate $S(t)$, a vector of network and ecological observables: population, fragmentation (1 minus largest-component share), degree- and betweenness-concentration, Freeman centralization [18], per-capita costly-helping rate, energy Gini, turnover, and mean traits (logged, never broadcast). The broadcast $b(t)$ carries five of these, normalized to $[0, 1]$: *fragmentation*, *degree concentration* (top-5% degree share, distinct from Freeman centralization), *cooperation*, *inequality* and *turnover*. Six conditions define what agents receive (Table 2). Distorted broadcasts come in three modes: *invert* ($b_k = 1 - s_k$ componentwise), *crisis* (fixed alarming values: fragmentation 0.9, degree concentration 0.9, cooperation 0.1, inequality 0.9, turnover 0.9), and *utopia* (reassuring: 0.05, 0.2, 0.9, 0.1, 0.2).

3.3 Response model

Each step, agent i selects one action a with probability proportional to

$$w_i(a, t) = w_a(\theta_i, x_i(t)) + \gamma_i g f_a(b(t)), \quad (1)$$

Table 1. Model parameters (defaults; campaign overrides noted in the text). Shock at $t_{\text{shock}} = 2000$ (1500 in the $2\times$ scale check).

initial population	150	resource capacity K	6000	observer interval Δt	10 steps
population cap	1200	base regeneration	60/step	feedback gain g	0.8 (swept 0.2–1.6)
initial edges/agent	3 (random)	regeneration rate r_{reg}	0.12	mutation σ	0.05
initial energy	20	harvest cap	2.5/step	initial traits	mean 0.5, s.d. 0.15
metabolism	0.8/step	connect cost / share	0.25 / 3.0	reprod. threshold / cost	30 / 16
run length	3000–4000 steps	min. reproduction age	15 steps	lifespan (mean $\pm$ s.d.)	400 ± 80 steps

Table 2. Experimental conditions; all start from matched parameter distributions with paired seeds.

Condition	Observer	Agents receive	Role
A — Local only	off	nothing	baseline
B — Observed, blind	on	nothing	measurement control
C — True feedback	on	accurate $S(t)$	main treatment
F — False feedback	on	distorted $S(t)$	reflexivity probe
N — Noise feedback	on	uniform random, matched ranges	stimulation control
R — Replayed self-model	on	another run’s genuine $S(t)$	structure control

where $x_i(t)$ is i ’s local state, g the global feedback gain, f_a the response rule, and the base propensities are harvest $1 + \mathbf{1}[e_i < 15]$ (yield scaled by θ_{risk}), share $0.6\theta_{\text{coop}}\theta_{\text{share}}$, connect $0.4\theta_{\text{soc}} + 0.3\theta_{\text{expl}}$, prune 0.08, reproduce $0.5\mathbf{1}[e_i \geq 30]$ and idle 0.2 ($e_i = \text{energy}$). In the primary *corrective* regime agents repair reported deficits: $f_{\text{connect}} = \max(0, b_{\text{frag}} - 0.3)$, $f_{\text{share}} = \max(0, 0.5 - b_{\text{coop}}) + 0.5\max(0, b_{\text{ineq}} - 0.5)$, $f_{\text{harvest}} = \max(0, b_{\text{turn}} - 0.5)$, and $f_{\text{prune}} = 0.3\max(0, b_{\text{conc}} - 0.6)$. In the *conformist* regime agents imitate the reported norm instead: $f_{\text{share}} = b_{\text{coop}}$, $f_{\text{prune}} = 0.5b_{\text{frag}}$, $f_{\text{connect}} = \max(0, 0.6 - b_{\text{frag}})$, and $f_{\text{harvest}} = b_{\text{ineq}}$. Two mechanisms act on *target selection* rather than action choice, in both regimes: while reported degree concentration exceeds 0.6 a pruning agent cuts its highest-degree neighbor (ties by id) with probability γ_i — g scales action propensities only, not this probability — else a uniform neighbor; and new links attach to a neighbor-of-a-neighbor with probability 0.7 when available, else a uniform non-neighbor — the triadic closure behind Sect. 3.4. All rule constants are fixed across experiments (Table 1). They are deliberately stylized — mechanism probes, not calibrated models of behavior. The abstraction is intentional: the model is designed to isolate causal pathways through which self-model feedback alters an adaptive system, not to estimate human behavioral parameters.

3.4 Complex-network structure

At $t = 2,000$ the giant component (six graphs: A and $C \times 3$ seeds) has mean path length $L \approx 1.8$ –2.5, clustering 1.8 – $2.9\times$ a matched Erdős–Rényi null — though against the stricter degree-preserving configuration-model null most of that excess traces to the degree sequence (ratios 1.0–1.6), so we make no small-world claim — and degree CV 0.6–0.74. Feedback acts on a small, dense, degree-heterogeneous clustered graph, not homogeneous mixing.

3.5 Reproducibility

The simulation is deterministic given (configuration, seed): one pseudorandom number generator (PRNG) drives behavioral choices, a dedicated stream generates broadcast signals (so constructing noise never advances the behavioral PRNG), agents act in sorted order, and a state hash certifies replay across machines.

4 Experimental Design

The standardized shock removes, in a single step at $t_{\text{shock}} = 2000$ (Table 1), the highest-degree living agents (ties broken by agent id): a fixed 10 hubs in the mild variant, the top $\max(1, \lfloor 0.4 |V_{\text{live}}| \rfloor)$ in the harsh variant. Four outcome *families* were fixed before the first campaign: (1) post-shock recovery time — the first $t > t_{\text{shock}}$ at which the living population regains 90% of its pre-shock baseline, defined as its mean over $[t_{\text{shock}} - 500, t_{\text{shock}})$; (2) mean post-shock fragmentation; (3) per-capita costly-helping rate, averaged after a 500-step burn-in; (4) self-model correspondence, operationalized as *story pull*: with $s_k(t)$ the value of macrostate component k at observer tick t , $b_k(t)$ the broadcast value, and $\varepsilon = 0.02$,

$$P = \left\langle \text{sign}(b_k(t) - s_k(t)) (s_k(t+1) - s_k(t)) \right\rangle_{(t,k) : |b_k(t) - s_k(t)| > \varepsilon}, \quad (2)$$

averaged over qualifying (tick, component) pairs across the five broadcast components; $P > 0$ means the macrostate moves toward the reference where the two differ. The threshold makes P well defined under false and noise broadcasts (undefined for C, whose broadcast equals the state).

Outcome pre-specification. Eq. 2 is a *post hoc* directional refinement of the fourth family (the original distance form is degenerate under truthful feedback; both give the same raw F-vs-N ordering). No pre-specified outcome was dropped; all other reported measures, including the evolved trait change, are exploratory.

Passive counterfactual. P can also be evaluated on trajectories that never received the reference: for each treated run we compute P on the paired-seed *untreated* trajectory (A/B) scored against a matched reference of the same type, defining the *counterfactual performativity gap* $\Delta P = P_{\text{actual}} - P_{\text{passive}}$ (Sect. 5.2). Campaigns: a *flagship* run (conditions A/B/C/F_{invert}/N $\times$ 50 paired seeds, mild shock), a *harsh-shock* replication (40% hub removal, $\times$ 30 seeds), a *sensitivity sweep* (feedback gain $g \in \{0.2, 0.8, 1.6\} \times \{A, C, F_{\text{invert}}, F_{\text{crisis}}, F_{\text{utopia}}, N\} \times$ 20 seeds), a $2\times$ *scale check* (doubled population and resources), a conformist *response-regime campaign* (15 seeds per cell), and a *replay control* (30 receiving runs on the harsh-shock seed list plus 15 auxiliary source runs): the core campaigns comprise $250 + 150 + 360 + 32 + 90 + 30 + 15 = 927$ runs of 3,000–4,000 steps; the three architectural controls (Sect. 5.5) add $300 + 300 + 400$ and the ecological ensemble (Sect. 5.7) 1,200, for 3,127 in total. Because conditions share

Table 3. Main paired comparisons (Wilcoxon signed-rank on per-seed differences; matched-pairs rank-biserial r ; bootstrap 95% CI of the median difference). FL = flagship (50 pairs), HS = harsh shock (30 pairs). “adj.” compares $\Delta P = P_{\text{actual}} - P_{\text{passive}}$ between conditions (Sect. 5.2). Multiplicity: over the eight rows carrying a point p , Benjamini–Hochberg and the stricter Holm–Bonferroni leave every decision unchanged; over all 80 paired contrasts in the paper, BH still changes nothing here and Holm demotes only the weakest row ($p = 9.5 \times 10^{-4}$), other demotions falling in exploratory checks, which are read through effect size and replication rather than adjusted significance.

Outcome	Comparison	Δ median [CI]	r	p
all state outcomes (FL, HS)	B vs A	all differences $\equiv 0$	—	—
story pull, raw (HS)	F vs N	+0.0066 [0.0037, 0.0095]	0.89	1.7×10^{-6}
story pull, adj. (HS)	F vs N	+0.0022 [−0.0003, 0.0034]	0.36	0.084
story pull, adj. (FL)	F vs N	+0.0004 [−0.0016, 0.0022]	0.08	0.63
cooperation (FL)	C vs A	+0.200 [0.172, 0.239]	1.00	8.9×10^{-15}
cooperation (HS)	C vs A	+0.209 [0.169, 0.248]	0.99	9.3×10^{-9}
fragmentation (FL)	F vs A	−0.017 [−0.034, −0.006]	−0.52	9.5×10^{-4}
$\Delta\bar{\gamma}$ (HS)	F vs A	−0.336 [−0.456, −0.281]	−1.00	1.9×10^{-9}
$\Delta\bar{\gamma}$ (HS)	N vs A	−0.323 [−0.431, −0.240]	−0.94	2.0×10^{-7}
recovery time (FL, HS)	C, F, N vs A	Δ median = 0 $ r < 0.5$		> 0.22

seed lists the design is paired; the primary analysis is a Wilcoxon signed-rank test on per-seed differences with matched-pairs rank-biserial r and a bootstrap 95% CI on the median difference. Unpaired Mann–Whitney U with Cliff’s δ is a secondary robustness view; all inferential claims follow the paired analysis. Per-run distributions for every reported outcome ship with the released data; cells carry 15–100 paired seeds; the smaller checks — $2\times$ scale ($n = 8$), the ecological ensemble (10 per ecology) and the structural nulls (six graphs) — are flagged where they appear.

5 Results

Table 3 summarizes the main paired comparisons; the following subsections unpack them.

5.1 Measurement alone is causally inert

Under paired seeds, condition B is trajectory-identical to A: every per-seed difference is exactly zero, on every outcome, in both campaigns. The instrument acquires causal power only when it broadcasts.

5.2 Apparent story pull is explained by a passive counterfactual

Except where stated, results concern the corrective regime; Sect. 5.4 shows how they change under conformity. Evaluated naively, the story-pull statistic appears to show strong self-fulfilling steering. Raw P (harsh shock; medians with bootstrap 95% CIs) is positive in all three non-truth conditions: F 0.0078

[0.0066, 0.0100] and N 0.0014 [0.0001, 0.0023]. It is largest for a sibling world’s replayed genuine self-model, condition R: 0.0241 [0.0197, 0.0265] (30 receivers on the harsh-shock seed list, one of 15 sources each). The hierarchy $R > F > N$ is reliable — F vs N $r = 0.89$, $p \approx 1.7 \times 10^{-6}$, surviving one receiver per source ($n = 15$, $r \geq 0.98$) and a source-clustered bootstrap.

A short observation shows why raw P can mislead. Under the inverted reference $b_t = 1 - s_t$ we have $\text{sign}(b_t - s_t) = \text{sign}(\frac{1}{2} - s_t)$. Suppose an untreated component simply mean-reverts toward mid-range: $s_{t+1} - s_t = -\kappa(s_t - \frac{1}{2}) + \eta_t$ with $\kappa > 0$ and $\mathbb{E}[\eta_t | s_t] = 0$. Then $\mathbb{E}[\text{sign}(b_t - s_t)(s_{t+1} - s_t) | s_t] = \kappa |s_t - \frac{1}{2}| > 0$: ordinary mean reversion masquerades as attraction toward the mirror, with zero causal feedback. Post-shock relaxation plays the same role for the replayed reference, which tracks where recovering trajectories go anyway. The remedy is the *counterfactual performativity gap* ΔP : score the paired-seed *untreated* trajectory (A, or the identical B) post hoc against a matched — not resampled — reference of the same type: for N and R the twin gets the *same realized* sequence its treated partner received, index-aligned on the observer grid; for invert, $b_t^* = 1 - s_t$ on the twin itself. Untreated worlds “move toward” these references *more* than treated worlds move toward their actual broadcasts: passive medians 0.0119, 0.0300, 0.0073. Every ΔP is therefore negative: -0.0036 (F), -0.0058 (R), -0.0055 (N), each $p < 2 \times 10^{-5}$. Between-condition contrasts of ΔP are significant in neither campaign (F vs N: harsh $p = 0.084$, flagship $p = 0.63$; R vs N: $n=30$, $p = 0.95$). Raw P is thus dominated by the intrinsic dynamics above, which align best with realistic references and worst with noise (excluding the shock transition shifts every median by < 0.0004). Response activation cannot explain the hierarchy either: R is the most quiescent of the three ($\langle |b_k - s_k| \rangle$ 0.10, against 0.64 F and 0.36 N) yet has the largest raw pull. Nor does the threshold decide it: under $\varepsilon \in \{0.01, 0.02, 0.05\}$ every ΔP stays negative (all $p < 2 \times 10^{-5}$) and the adjusted F-vs-N contrast stays ns ($p \geq 0.052$). Observer cadence is not a comparable robustness axis: retiming the observer retimes the feedback dose itself, so Δt is a treatment parameter; supplementary cadence-resolved analyses accompany the released data.

Under the primary corrective regime Eq. 2 thus provides *no evidence of aggregate content-specific attraction beyond intrinsic dynamics*. The small negative ΔP , similar across signal types, is consistent with broadcasts mildly dampening that drift. The methodological implication is that a statistic natural in empirical performativity studies can register large, internally consistent but entirely artifactual “steering” unless checked against passive twins — we offer ΔP as a paired diagnostic for designs with an untreated twin, not a general identification theorem.

Structural spillovers, by contrast, are real. Agents persistently told the network is fragmenting over-connect it: mean degree 19.8 vs 12.3 in matched flagship runs at $t = 2000$ (paired $\Delta \text{median} +7.2$, $p \approx 7 \times 10^{-10}$). Post-shock fragmentation falls with it (flagship paired $\Delta \text{median} -0.017$, $r = -0.52$, $p \approx 10^{-3}$; directionally consistent but ns under harsh shock). False and noise broadcasts densify the network, and also flatten it. In the baseline ecology, relative to A




Fig. 2. Sensitivity sweep (40% hub removal; 20 seeds per cell; medians). Left: attention varies with broadcast type and gain; the flattering (utopia) lie tracks the no-feedback baseline exactly (dashed over grey). Middle: under correction the alarmist (crisis) lie does *not* come true. Right: deficit signals mobilize costly helping.

(paired, harsh shock, per-run medians over $t \geq 500$), mean degree rises +1.1 under truth ($r = 0.49$), +3.8 under the inverted lie ($r = 0.80$) and +2.8 under noise ($r = 0.72$). On the same runs and window, Freeman centralization falls under both distortions (F: -0.026 , $r = -0.74$; N: -0.010), as does betweenness concentration (F: -0.007 , $r = -0.67$; N: -0.008).

5.3 The type of lie sets what changes

The three distortion modes produce qualitatively different *outcome* regimes (Fig. 2). The *inverted* self-model, always reporting the mirror of reality, keeps four of five corrective responses active on $\geq 90\%$ of ticks: it mobilizes costly helping (+0.106 over baseline, flagship $r = 0.82$) and drives the over-connection above. The *crisis* self-model is structurally self-defeating — corrections push fragmentation away from the alarmist report — yet its permanent cooperation-deficit signal mobilizes the highest costly helping of any lie under correction (0.35–0.45 across gains; only truth at $g = 1.6$ exceeds them, 0.48). The *utopia* self-model is behaviorally inert: responses fire only on reported deficits, so at every gain flattery is indistinguishable from disconnecting the channel — under purely corrective rules, flattery equals silence. Caveat: these response *directions* follow from Eq. 1 by construction; the magnitudes, spillovers and evolutionary consequences below do not.

5.4 The response regime decides which lies come true

Conformist agents reverse the taxonomy exactly as the mechanism predicts (Table 4). The *crisis* lie, self-defeating under correction, becomes strongly self-fulfilling under conformity: told the world is fragmenting, imitators cut ties and fragmentation rises an order of magnitude above baseline (0.400 vs 0.038). The *utopia* lie flips from behaviorally inert to benevolently self-fulfilling: “everyone cooperates” becomes true by imitation (cooperation 0.428). Even noise mobilizes: imitating randomly reported cooperation levels raises costly helping to 0.445, as high as any lie, since imitation amplifies whatever is reported. Which false self-models come true — in the outcome sense — is therefore a property of the lie-response pair, not the lie alone: with macrostate update $S_{t+1} = \Phi(S_t, \mathcal{R}(\mathcal{B}(S_t)))$, self-fulfillment is a property of the composition $\mathcal{R} \circ \mathcal{B}$, not of the broadcast map $\mathcal{B}$ alone.

Table 4. The same lie under two response regimes (medians over the 15 seeds shared by both campaigns; harsh shock, $g = 0.8$). Baseline (A), identical in both regimes: cooperation 0.183, fragmentation 0.038.

Lie	Corrective		Conformist	
	coop.	fragm.	coop.	fragm.
Invert	0.256	0.020	0.419	0.054
Crisis	0.414	0.018	0.255	0.400
Utopia	0.183 (inert)	0.038 (inert)	0.428	0.025

5.5 Attention declines under consequential broadcasts

Mean heritable global sensitivity $\bar{\gamma}$ changes over ~ 20 generations of lineage depth as an emergent, unprogrammed function of broadcast quality. At $g = 0.8$ (harsh shock): no broadcast -0.03 (neutral drift), truth -0.10 , noise -0.34 ($r = -0.94$, $p \approx 2 \times 10^{-7}$), false -0.42 (Δ median -0.336 $[-0.456, -0.281]$, $r = -1.00$, $p \approx 1.9 \times 10^{-9}$). The sweep is dose-ordered: the decline under the inverted broadcast deepens from -0.12 ($g = 0.2$) to -0.45 ($g = 1.6$), and at the highest gain even *accurate* feedback declines markedly (-0.24). The categorical ordering compresses into one continuous relationship: evolved attention tracks the broadcast-implied corrective-drive intensity $I = g \langle \sum_a f_a(b(t)) \rangle$ — computable from the broadcast alone — across the 15 sweep cells almost perfectly (Spearman $\rho = -0.99$; $\rho = -0.83$ at run level, $n = 300$). Utopia ($I = 0$) is untouched, truth declines less where its corrective drive is smaller at equilibrium, and high-drive broadcasts pay most: the response appears organized more by corrective-drive intensity than by veracity alone.

Three caveats. I omits the per-agent factor γ_i , the quantity under selection, and comes from realized rather than independently randomized trajectories; we treat it as descriptive, not causal mediation. It is compared only within corrective-regime sweeps — a within-regime ordering, not normalized across regimes. And the analysis is population-level: whether the $\bar{\gamma}$ shift arises from within-lineage movement, differential reproduction, or compositional change requires agent-level lineage tracking, left to future work.

Not an architectural artifact. Two architectural asymmetries could in principle manufacture this pattern. First, feedback adds weight only to non-reproductive actions. Write $D(b) = \sum_a f_a(b)$ for that added mass, whose gain-scaled mean is I , and let W_0 be the base total weight. Then $P(\text{reproduce} | \gamma_i, b) = w_r / (W_0 + \gamma_i g D(b))$ falls with γ_i whenever $D(b) > 0$. Second, hub-targeted pruning fires with probability γ_i rather than $\gamma_i g$, a second, differently scaled channel. We reran all five broadcast conditions against each, 300 runs apiece. The reproduction-neutral variant rescales w_r by $(1 + \gamma_i g D/N_0)$ (N_0 the base non-reproductive mass), holding selection probability at its no-broadcast value, other feedback weights unchanged; the second fixes hub-pruning probability at 0.5 for every agent.

The pattern survives both controls: over cells with a negative baseline $\Delta\bar{\gamma}$, median retention is 98% (reproduction-neutral) and 100% (γ -free pruning). The

intensity correlation is almost unchanged ($\rho = -0.98$ and -1.00). Under the reproduction control, truth in fact declines *more* at $g \geq 0.8$. One γ -free-pruning cell of fifteen shifts uncorrected (noise at $g = 0.8$, attenuating to 77% of baseline, paired $p = 0.033$); the intensity ordering tightens. Selection on γ is ecological rather than an artifact of either channel, though neither control neutralizes the energy cost of induced actions; a cost-neutral control separating costly responsiveness from unreliable information is future work. The decline is not specific to the corrective rule form: under the conformist mapping $\bar{\gamma}$ still falls in every broadcast condition (paired $\Delta_{\text{median}} -0.14$ to -0.57 , all raw $p \leq 0.022$), though neither that ordering nor I carries over ($\rho = +0.10$, $n = 5$), as the scope above implies.

Evolvable response strength. With f_a fixed, lowering γ is the population’s only route to disengagement. Making the polarity and strength $\rho_a \in [-2, 2]$ of each action channel heritable — $\rho_a = 1$ reproduces the corrective rule bit-identically, 0 ignores that channel, <0 reverses it — opens a second route, and $\bar{\gamma}$ still declines (vs. A, $n = 100$: C -0.125 , F -0.222 , N -0.212 ; all raw $p < 10^{-4}$). On matched seeds this attenuates the decline under false and noise feedback ($+0.126$, $r = 0.66$; $+0.165$, $r = 0.65$; both raw $p \approx 0.001$) but not under truth ($p = 0.90$): populations use both routes exactly where disengagement pays. The evolutionary result is therefore not an artifact of making γ the sole adjustable response; the individual-level pathway remains open.

5.6 Truthful feedback doubles cooperation; main-campaign recovery is unchanged

True feedback roughly doubles costly helping (flagship: 0.397 vs 0.179; paired $\Delta_{\text{median}} +0.200$ [0.172, 0.239], $r = 1.00$, $p \approx 9 \times 10^{-15}$; replicated under harsh shock, $r = 0.99$), rising monotonically with gain ($0.27 \rightarrow 0.48$): deficits reported at the population level are repaired by uncoordinated individual action. By contrast, a pre-specified null: recovery time showed no feedback effect in either main campaign, in any condition, even with 40% of hubs removed. Every paired Δ_{median} is 0 and every $p > 0.22$ (baseline medians 10 and 40 steps, near the $\Delta t = 10$ observer resolution floor). One lead runs the other way: at $2 \times$ scale the inverted lie delayed recovery in every non-tied pair ($+10$ steps, $r = 1.00$; five pairs, smallest attainable $p = 0.0625$) and does not replicate at base scale. Under this recovery criterion, resilience appears predominantly demographic rather than informational.

5.7 Robustness to system scale and to the baseline ecology

Re-running A, C, F_{invert} , and N at doubled population and resources (steady-state ≈ 110 – 130 vs ≈ 60 ; $n = 8$) preserves the main effects: cooperation under truth 0.391 vs 0.173 ($\Delta_{\text{median}} +0.151$ [0.119, 0.247], $r = 1.00$, $p = 0.008$) and the attention ordering (F vs A -0.299 [$-0.388, -0.212$], $r = -1.00$, $p = 0.008$; N vs A -0.219 , $r = -0.89$, $p = 0.023$).

Ecological robustness. Perturbing six ecological parameters by $\pm 25\%$ gives 24 alternative ecologies, all passing viability criteria fixed on the *no-broadcast* condition before any treatment ran. Across $A/B/C/F/N \times 10$ paired seeds each (1,200 runs): cooperation rose under truth in 24/24, evolved attention fell under the inverted lie and noise in 24/24, measurement without broadcast stayed exactly inert in every paired run, and $F \leq N \leq C \leq A$ held in 21/24. The recovery null persisted; truthful densification did not (16/24 correct-signed), motivating Sect. 5.2’s narrower claim. Read as replication, this is *local* ecological robustness, not arbitrary-ecology generality.

6 Discussion and Limitations

First-order response directions are encoded in the rules; our claims concern the macroscopic, structural, and evolutionary consequences. First, the methodological lesson: a statistic that looked like compelling evidence — positive, replicated, largest for realistic signals — was fully explained by intrinsic dynamics on passive twins; empirical tests face this identification problem [22] typically without the untreated twin our platform provides. Second, the substantive findings: broadcasts change behavior and structure with a regime-dependent taxonomy of lies; more importantly, reliance on the collective self-model becomes an evolving population property. The response mappings are controlled mechanism probes, not calibrated behavioral models: they isolate how receiver transformations turn identical broadcasts into different outcomes; their strength and polarity are evolvable (Sect. 5.5) but the basis stays mechanistic; whether fully learned policies preserve the taxonomy is open. Further limitations: local rather than arbitrary-ecology generality (Sect. 5.7), and a single adaptive population network per run — competing populations or coupled self-model ecosystems remain open.

Metric-driven infrastructures — autoscalers, observability pipelines, ranking systems [4] — instantiate the same feedback loop: under deficit-driven responses *flattering telemetry can be operationally equivalent to no telemetry* while raising no alarms; agent populations sharing dashboards [15] are a natural empirical target, given a passive baseline before alignment is read as steering.

7 Conclusion

We introduced a reproducible causal laboratory for reflexive self-modeling in adaptive networks and showed that (i) measurement without broadcast is trajectory-identical to baseline, while broadcasting changes behavior, structure, and evolution; (ii) apparent “story pull” vanishes against a passive counterfactual, a caution for empirical performativity claims; (iii) which lies come true, in the outcome sense, is set by the lie–response pair, not the lie alone; and (iv) heritable attention to the collective self-model declines across generations, tracking broadcast-implied corrective drive (Sect. 5.5). Future work: learned response policies and language-model agents.

Reproducibility. Code, configurations, analysis scripts, pinned dependencies, and deterministic regeneration of every run: github.com/Ivan825/Project-ONE, tag `cn2026-camera-ready`.

References

1. Gross, T., Blasius, B.: Adaptive coevolutionary networks: a review. *J. R. Soc. Interface* 5, 259–271 (2008)
2. Holme, P., Saramäki, J.: Temporal networks. *Phys. Rep.* 519, 97–125 (2012)
3. Soros, G.: Fallibility, reflexivity, and the human uncertainty principle. *J. Econ. Methodol.* 20, 309–329 (2013)
4. Espeland, W.N., Sauder, M.: Rankings and reactivity: how public measures recreate social worlds. *Am. J. Sociol.* 113, 1–40 (2007)
5. Merton, R.K.: The self-fulfilling prophecy. *Antioch Rev.* 8, 193–210 (1948)
6. Strathern, M.: ‘Improving ratings’: audit in the British university system. *Eur. Rev.* 5, 305–321 (1997)
7. Albert, R., Jeong, H., Barabási, A.-L.: Error and attack tolerance of complex networks. *Nature* 406, 378–382 (2000)
8. Kuramoto, Y.: *Chemical Oscillations, Waves, and Turbulence*. Springer (1984)
9. Skyrms, B.: *Signals: Evolution, Learning, and Information*. Oxford Univ. Press (2010)
10. Epstein, J.M., Axtell, R.: *Growing Artificial Societies*. Brookings/MIT Press (1996)
11. Bonabeau, E.: Agent-based modeling: methods and techniques for simulating human systems. *PNAS* 99, 7280–7287 (2002)
12. Kazil, J., Masad, D., Crooks, A.: Utilizing Python for agent-based modeling: the Mesa framework. In: *SBP-BRiMS*. Springer (2020)
13. Hagberg, A., Schult, D., Swart, P.: Exploring network structure, dynamics, and function using NetworkX. In: *SciPy* (2008)
14. Foerster, J., Assael, Y.M., de Freitas, N., Whiteson, S.: Learning to communicate with deep multi-agent reinforcement learning. In: *NeurIPS* (2016)
15. Park, J.S., O’Brien, J., Cai, C.J., Morris, M.R., Liang, P., Bernstein, M.S.: Generative agents: interactive simulacra of human behavior. In: *UIST* (2023)
16. Castellano, C., Fortunato, S., Loreto, V.: Statistical physics of social dynamics. *Rev. Mod. Phys.* 81, 591–646 (2009)
17. Perc, M., Jordan, J.J., Rand, D.G., Wang, Z., Boccaletti, S., Szolnoki, A.: Statistical physics of human cooperation. *Phys. Rep.* 687, 1–51 (2017)
18. Freeman, L.C.: Centrality in social networks: conceptual clarification. *Soc. Netw.* 1, 215–239 (1978)
19. Perdomo, J.C., Zrnic, T., Mendler-Dünner, C., Hardt, M.: Performative prediction. In: *ICML. PMLR* 119 (2020)
20. Wang, X., Yau, C.-Y., Wai, H.-T.: Network effects in performative prediction games. In: *ICML. PMLR* 202 (2023)
21. Góis, A., Mofakhami, M., Santos, F.P., Gidel, G., Lacoste-Julien, S.: Performative prediction on games and mechanism design. In: *AISTATS. PMLR* 258, 1855–1863 (2025)
22. Cheng, G., Hardt, M., Mendler-Dünner, C.: Causal inference out of control: estimating performativity without randomization. In: *ICML. PMLR* 235, 8077–8103 (2024)
23. Hardt, M., Mendler-Dünner, C.: Performative prediction: past and future. *Stat. Sci.* 40, 417–436 (2025)
24. Iwase, T., Tadokoro, Y., Fukuda, D.: Self-fulfilling signal of an endogenous state in network congestion games. *Netw. Spat. Econ.* 17, 889–909 (2017)